a KVA projector

One middle state feeds a trained head per late layer; the heads write keys and values directly.

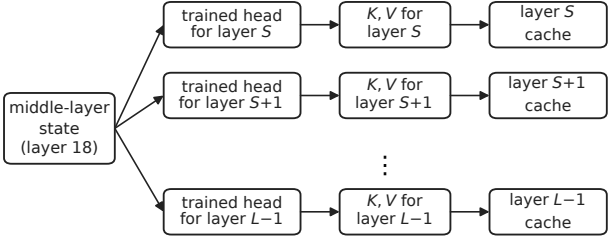

Each layer's own key and value projections are bypassed.

b RidgeFill

One split state feeds a closed-form map per late layer; each layer runs its own projections on the predicted input.

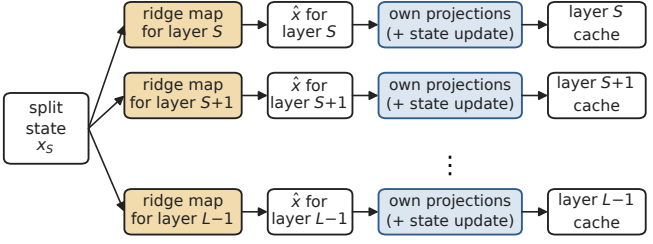

No layer waits on another. Exact rows (§2.4) skip the maps.