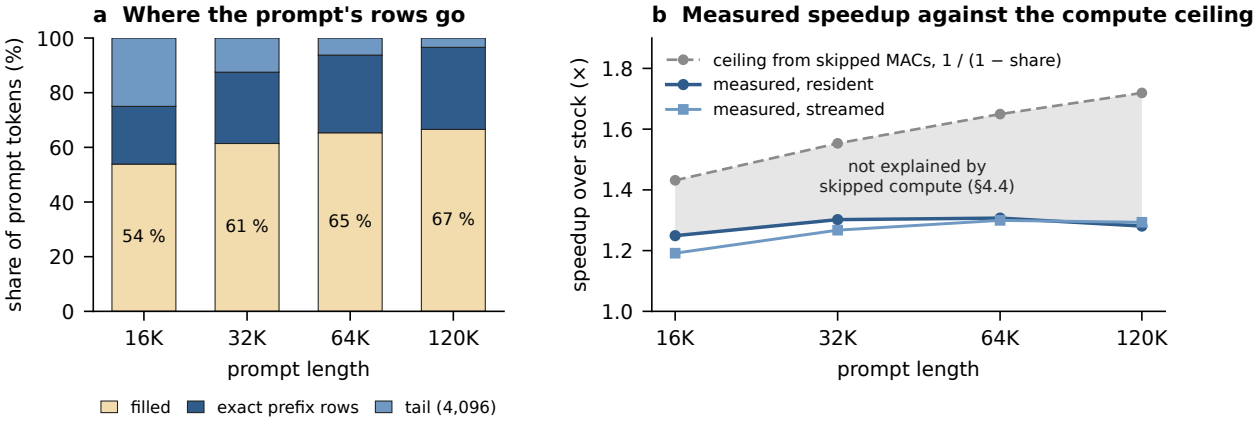
a Where the prompt's rows go
share of prompt tokens (%)
100
80
60
40
20
0
54 %
61 %
65 %
67 %
16K
32K
64K
120K
prompt length
filled
exact prefix rows
tail (4,096)
b Measured speedup against the compute ceiling
speedup over stock (×)
1.8
1.6
1.4
1.2
1.0
ceiling from skipped MACs, 1 / (1 − share)
measured, resident
measured, streamed
not explained by
skipped compute (§4.4)
16K
32K
64K
120K
prompt length