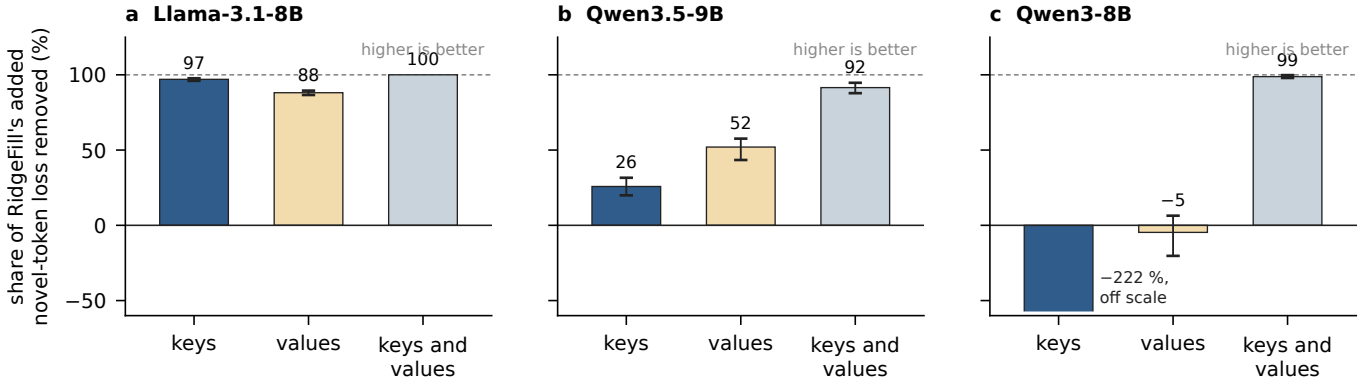
a Llama-3.1-8B
b Qwen3.5-9B
c Qwen3-8B
share of RidgeFill's added novel-token loss removed (%)
higher is better
100
50
0
−50
97
88
100
26
52
92
−5
99
−222 %, off scale
keys
values
keys and values
Exact counterpart spliced into every filled attention layer